

Øving 7

2022-05-30

Problem 1

- a) Provide a detailed explanation of the algorithm that is used to fit a regression tree. What is different for a classification tree?

We divide the predictor space into J distinct and non-overlapping regions R_1, R_2, \dots, R_J . For every observation that falls into the region R_j , we make the same prediction, which is simply the mean of the response values for the training observations in R_j . The goal is to find R_j such that the RSS is minimized.

$$RSS = \sum_{j=1}^J \sum_{i \in R_j} (y_i - \hat{y}_{R_j})^2$$

Where \hat{y}_{R_j} is the response for the training observations within the j 'th box.

For classification trees we predict a qualitative response rather than a quantitative one. We predict that each observation belongs to the most commonly occurring class of training observations in the region to which it belongs. We can NOT use RSS for binary splits.

- b) What are the advantages and disadvantages of regression and classification trees?

Advantages:

- Simple and useful for interpretation.
- Easier to explain than regression.
- More closely mirrors human decision making.
- Bagging, random forests and boosting methods grow multiple trees, which results in improvements for prediction accuracy.
- Displayed graphically and therefore interpretable.
- Trees can easily handle qualitative predictors without the need to create dummy variables.

Disadvantages:

- Bagging, random forests and boosting methods grow gives loss in interpretation.
- Not competitive with the best supervised learning methods in terms of prediction accuracy.
- Trees generally do not have the same level of predictive accuracy as other regression and classification methods.

- c) What is the idea behind bagging and what is the role of the bootstrap? How do random forests improve that idea?

Bootstrap aggregation, or bagging is a procedure for reducing the variance of a statistical learning method.

Averaging a set of observations reduces the variance by a factor $1/n$. This is not practical because we generally don't have access to multiple training sets. Instead, we can bootstrap.

We generate B different bootstrapped training sets. We then train our method on the b th bootstrapped training set in order to get the estimated $\hat{f}_b(x)$, the prediction at point x . We then average all the predictors.

Random forests provide an improvement over bagged trees by way of a small tweak that decorrelated the trees. This reduces the variance when we average the trees.

As in bagging, we build a number of decision trees on bootstrapped training samples.

When building these decision trees, each time a split in a tree is considered, a random selection of m predictors is chosen as split candidates from the full set of p predictors. The split is allowed to use only one of those m predictors.

A fresh selection of m predictors is taken at each split, and typically we choose $m = \sqrt{p}$. That is the number of predictors considered at each split.

- d)
- e)

Problem 2

- a) Split the data into a training and a test set,

```
library(ISLR)

data("Carseats")

set.seed(4268)

n = nrow(Carseats)

train = sample(1:n, 0.7 * nrow(Carseats), replace = F)
test = (1:n)[-train]

Carseats.train = Carseats[train,]
Carseats.test = Carseats[-train,]

head(Carseats)
```

```
##   Sales CompPrice Income Advertising Population Price ShelfLoc Age Education
## 1  9.50      138     73          11         276   120      Bad  42         17
## 2 11.22      111     48          16         260    83     Good  65         10
## 3 10.06      113     35          10         269    80   Medium  59         12
## 4  7.40      117    100           4         466    97   Medium  55         14
## 5  4.15      141     64           3         340   128     Bad  38         13
## 6 10.81      124    113          13         501    72     Bad  78         16
##   Urban  US
## 1   Yes Yes
## 2   Yes Yes
## 3   Yes Yes
## 4   Yes Yes
## 5   Yes  No
## 6    No Yes
```

- b) Fit a regression tree to the training set. Plot the tree, and interpret the results. What test MSE do you obtain?

- c) Use cross-validation in order to determine the optimal level of tree complexity. Does pruning the tree improve the test MSE?

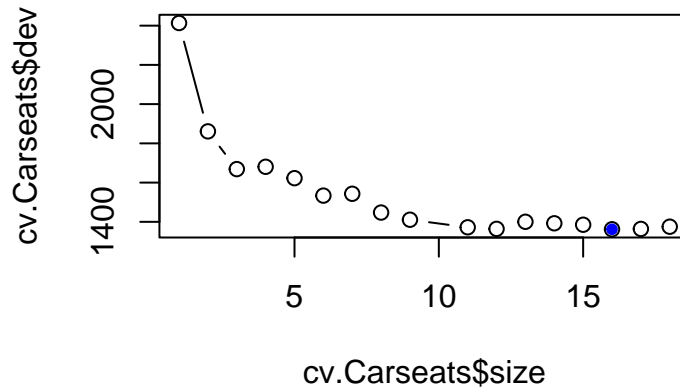
```
set.seed(4268)

cv.Carseats = cv.tree(tree.mod, FUN = prune.tree, K = 10)

tree.min = which.min(cv.Carseats$dev)

best = cv.Carseats$size[tree.min]

plot(cv.Carseats$size, cv.Carseats$dev, type = "b")
points(best, cv.Carseats$dev[tree.min], col = "blue", pch = 20)
```



We observe that the trees of the size 11 and 12 have similar results as the 16th. We might choose 11 for a simpler tree.

```
pruned_tree = prune.tree(tree.mod, best = 11)
pruned_preds = predict(pruned_tree, newdata = Carseats.test)

pruned_MSE = mean((pruned_preds - Carseats.test$Sales)^2)
pruned_MSE
```

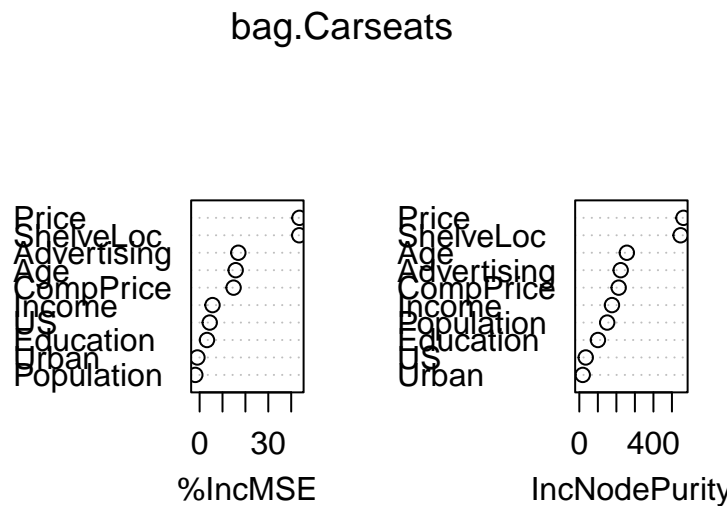
```
## [1] 4.378499
```

Which is a lower test MSE than the first tree. We plot the tree to view it.

```
plot(pruned_tree)
text(pruned_tree)
```



```
varImpPlot(bag.Carseats)
```



We observe that Shelveloc, Price, age and advertising are the most important variables.

- e) Use random forests to analyze the data. Include 500 trees and select 3 variables for each split. What test MSE do you obtain? Use the importance() function to determine which variables are most important. Describe the effect of m , the number of variables considered at each split, on the error rate obtained.

```
rf.Carseats = randomForest(Sales ~ ., data = Carseats.train, mtry = 3, ntree = 500, importance = TRUE)

yhat.rf = predict(rf.Carseats, newdata= Carseats.test)

rf_MSE = mean((yhat.rf - Carseats.test$Sales)^2)
rf_MSE
```

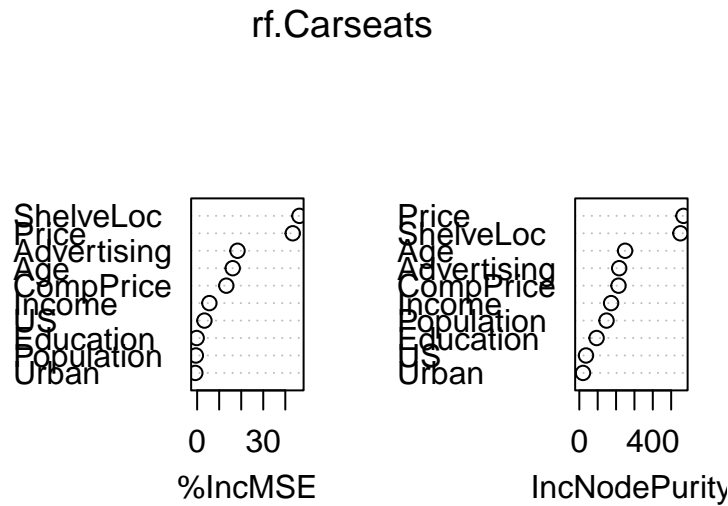
```
## [1] 2.25897
```

We use $p/m = 10/3$ trees (about 3 trees). We get an MSE a little larger than Bagging.

```
importance(rf.Carseats)
```

```
##           %IncMSE IncNodePurity
## CompPrice  13.235917    213.40899
## Income      5.5487005    173.08169
## Advertising 18.3241456    217.00013
## Population -0.5880682    147.97481
## Price      43.3832973    566.91128
## ShelfLoc   46.4146575    550.24111
## Age        16.1467391    249.21306
## Education  -0.2206834     94.00016
## Urban      -0.8309524     19.95026
## US         3.3052482     36.29111
```

```
varImpPlot(rf.Carseats)
```



- f) Finally use boosting with 500 trees, an interaction depth $d = 4$ and a shrinkage factor $\lambda = 0.1$, which is the default for the `gbm()` function. Compare MSE to the other methods.

```
#install.packages("gbm")
library(gbm)

r.boost = gbm(Sales ~ ., data = Carseats.train, distribution = "gaussian", n.tree = 500, interaction.depth = 4, shrinkage = 0.1)

boost_preds = predict(r.boost, newdata = Carseats.test)

boost_MSE = mean((boost_preds - Carseats.test$Sales)^2)
boost_MSE
```

```
## [1] 1.929048
```

We observe that the MSE is even lower than for the other methods.

- g) What is the effect of the number of trees (`ntree`) on the test error. Plot the test MSE as a function of `ntree` for both the bagging and the random forest method.

Problem 3

- b) Create a training set and a test set for the dataset.

```
#install.packages("kernlab")

library(kernlab)
data(spam)
n = nrow(spam)
```

```
train_index = sample(1:n, 0.7 * nrow(spam), replace = F)

spam.train = spam[train_index,]
spam.test = spam[-train_index,]
```

- c) Fit a tree to the training data with type as the response and the rest of the variables as predictors. Study the results by using the summary() function. Also create a plot of the tree. How many terminal nodes does it have?

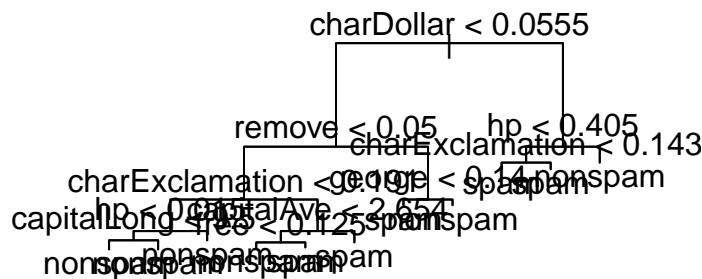
```
library(tree)

tree_model = tree(type ~ ., data= spam.train)
summary(tree_model)

##
## Classification tree:
## tree(formula = type ~ ., data = spam.train)
## Variables actually used in tree construction:
## [1] "charDollar"      "remove"          "charExclamation" "hp"
## [5] "capitalLong"     "capitalAve"      "free"            "george"
## Number of terminal nodes:  11
## Residual mean deviance:  0.5184 = 1664 / 3209
## Misclassification error rate: 0.09317 = 300 / 3220
```

We have 13 terminal nodes.

```
plot(tree_model)
text(tree_model)
```



- d) Predict the response on the test data. What is the misclassification rate?


```
tree_preds = predict(tree_model, newdata = spam.test, type = "class")  
missclass_table = table(tree_preds, spam.test$type)  
error_rate = 1-sum(diag(missclass_table))/sum(missclass_table)  
error_rate
```

```
## [1] 0.0948588
```

e)